

Aluno: Weld Lucas Cunha

Esse trabalho tem como base o código fornecido em aula pelo professor Samuel.

1. Set up

Imports

```
In [1]: from copy import deepcopy
        from typing import Tuple

        import numpy as np
        from numpy import ndarray
        import matplotlib.pyplot as plt
        import seaborn as sns
        import pandas as pd

        from sklearn.base import BaseEstimator, ClassifierMixin
        from sklearn.datasets import make_blobs
        from sklearn.metrics import classification_report, plot_confusion_matrix
        from sklearn.model_selection import train_test_split, RandomizedSearchCV
        from sklearn.preprocessing import MinMaxScaler
```

Creating fake data

```
In [2]: # fake data for testing
        X, y = make_blobs(n_samples=1000, n_features=2, centers=4, cluster_std=1.5, random_state=42)

        print(X.shape)
        print(y.shape)
        print(f'Labels: {np.unique(y)}')

        (1000, 2)
        (1000,)
        Labels: [0 1 2 3]
```

```
In [3]: # splitting into train and test
        X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

```
print(f'X_train.shape = {X_train.shape}')  
print(f'y_train.shape = {y_train.shape}')  
print(f'X_test.shape = {X_test.shape}')  
print(f'y_test.shape = {y_test.shape}')
```

```
X_train.shape = (800, 2)  
y_train.shape = (800,)  
X_test.shape = (200, 2)  
y_test.shape = (200,)
```


In [4]:

```
print(f'X_train.shape = {X_train.shape}')  
print(f'y_train.shape = {y_train.shape}')  
  
print(f'X_test.shape = {X_test.shape}')  
print(f'y_test.shape = {y_test.shape}')
```


```
X_train.shape = (800, 2)  
y_train.shape = (800,)  
X_test.shape = (200, 2)  
y_test.shape = (200,)
```

In [5]:

```
def plot_scatter(X_array, y_array, title=''):  
    plt.title(title)  
    sns.scatterplot(x=X_array[:, 0], y=X_array[:, 1], hue=y_array)  
    plt.xlabel('x1')  
    plt.xlabel('x2')  
    plt.show()  
  
plot_scatter(X_train, y_train, title='Training Samples')
```



```
In [6]: plot_scatter(X_test, y_test, title='Test Samples')
```



```
In [ ]:
```

2. Implementation

In [7]:

```

class LogisticRegression(ClassifierMixin, BaseEstimator):
    """Our Logistic Regression implemented from scratch."""

    def __init__(self, learning_rate : float = 0.001,
                  n_epochs : int = 1000, alpha: float = 0.0001, random_state : int = 42):
        """
        Parameters
        -----
        learning_rate : float, default=0.001
            Learning rate.
        n_epochs : int, default=1000
            Number of epochs for training (convergence stop).
        alpha : float, default=0.0001
            Constant that multiplies the regularization term.
            Use 0 to ignore regularization (standard Logistic Regression).
        random_state : int, default=42
            Seed used for generating random numbers.
        """

        assert (learning_rate is not None) and (learning_rate > 0.0), \
            f'Learning rate must be > 0. Passed: {learning_rate}'

        assert (n_epochs is not None) and (n_epochs > 0), \
            f'Number of epochs must be > 0. Passed: {n_epochs}'

        assert (alpha is not None) and (alpha >= 0), \
            f'Alpha should be >= 0. Passed: {alpha}'

        # public ==> todo mundo tem acesso para leitura e escrita direta
        # private ==> apenas a classe tem acesso para leitura e escrita direta
        self.learning_rate = learning_rate
        self.n_epochs = n_epochs
        self.alpha = alpha
        self.random_state = random_state

        # parameters to be trained/learned
        self.classes_ = []
        self.__indexes_dict = {}
        self.__w = None # weight array
        self.__b = None # bias
        self.__w_dict = {}

```

```

self.__b_dict = {}

# a special method used to represent a class object as a string, called with print() or str()
def __str__(self):
    msg = f'Learning rate: {self.learning_rate}\n' \
          f'Number of epochs: {self.n_epochs}\n' \
          f'Regularization constant (alpha): {self.alpha}\n' \
          f'Random state: {self.random_state}\n\n' \
          f'Trained?: {self.is_fitted()}\n'
    return msg

# getter: access the function as an attribute - it is not possible to set values through it
@property
def coef_(self) -> ndarray:
    """Return the weight matrix (learned parameters) if the estimator was fitted/trained.
    Otherwise, raise an exception.
    """
    assert self.is_fitted(), 'The instance is not fitted yet.'
    return self.__w_dict

# getter: access the function as an attribute - it is not possible to set values through it
@property
def intercept_(self) -> float:
    """Return the bias (learned intercept) if the estimator was fitted/trained.
    Otherwise, raise an exception.
    """
    assert self.is_fitted(), 'The instance is not fitted yet.'
    return self.__b

def is_fitted(self) -> bool:
    return self.__w is not None

def __sigmoid(self, z: ndarray) -> ndarray:
    return 1 / (1 + np.e ** (-z))

def __log_loss(self, y: ndarray, p_hat: ndarray, eps: float = 1e-15):
    '''Return the log loss for a given estimation and ground-truth (true labels).

```

log is undefined for 0. Consequently, the log loss is undefined for `p_hat=0` (because of `log(p_hat)`) and `p_hat=1` (because of `log(1 - p_hat)`). To overcome that, we clipped the probabilities to `max(eps, min(1 - eps, p_hat))`, where `eps` is a tiny constant.

Parameters

```
y : ndarray, shape (n_samples,)
    True labels of input samples.
p_hat : ndarray
    Estimated probabilities of input samples.
eps : float, default=1e-15
    Epsilon term used to avoid undefined log loss at 0 and 1.
```

Returns

```
log_loss : float
    Computed log loss.
...
```

```
p_hat_eps = np.maximum(eps, np.minimum(1 - eps, p_hat))

# shape: (n_samples,)
losses = -(y * np.log(p_hat_eps) + (1 - y) * np.log(1 - p_hat_eps))
log_loss = losses.mean()

return log_loss
```

```
def __gradient(self, X: ndarray, y: ndarray, p_hat: ndarray,
                w: ndarray, alpha: float) -> Tuple[ndarray, float]:
    '''Compute the gradient vector for the log loss with regards to the weights and bias.
```

Parameters

```
X: ndarray of shape (n_samples, n_features)
    Training data.
y: ndarray of shape (n_samples,)
    Target (true) labels.
p_hat : ndarray, shape (n_samples,)
    Estimated probabilities.
w : ndarray, shape (n_features,)
    Weight array.
alpha : float
    Regularization constant.
```

Returns

```

-----
Tuple[ndarray, float]:
    Tuple with:
        - a numpy array of shape (n_features,) containing the partial derivatives w.r.t. the weights; and
        - a float representing the partial derivative w.r.t. the bias.
    ...

# X.shape: (n_samples, n_features)
# y.shape == p_hat.shape: (n_samples,)
n_samples = X.shape[0]

regularization = alpha * w

error = p_hat - y # shape (n_samples,)
grad_w = (np.dot(error, X) / n_samples) + regularization # shape (n_features,)
grad_b = error.mean() # float

return grad_w, grad_b

def fit(self, X: ndarray, y: ndarray, verbose: int = 0):
    for i, val in enumerate(np.unique(y)):
        self.classes_.append(val)
        self.__indexes_dict[val] = i
        y_copy = np.zeros(y.shape)
        y_copy[y == val] = 1
        self.__fit(X, y_copy, verbose=verbose)
        self.__w_dict[i] = self.__w
        self.__b_dict[i] = self.__b

def __fit(self, X: ndarray, y: ndarray, verbose: int = 0):
    '''Train a Logistic Regression classifier.
    Parameters
    -----
    X: ndarray of shape (n_samples, n_features)
        Training data.
    y: ndarray of shape (n_samples,).
        Target (true) labels.
    verbose: int, default=0
        Verbose flag. Print training information every `verbose` iterations.

    Returns
    -----
    self : object
        Returns self.
    ...

```

```

#### CHECK INPUT ARRAY DIMENSIONS
assert X.ndim == 2, f'X must be 2D. Passed: {X.ndim}'
assert y.ndim == 1, f'y must be 1D. Passed: {y.ndim}'
assert X.shape[0] == y.shape[0], \
    f'X.shape[0] should be equal to y.shape[0], instead: {X.shape[0]} != {y.shape[0]}'
# alternatively
# X, y = check_X_y(X, y)

#### SETTING SEED
np.random.seed(self.random_state)

n_samples, n_features = X.shape

#### PARAMETER INITIALIZATION
# return values from the "standard normal" distribution.
w = np.random.randn(n_features) # shape: (n_features,)
b = 0.0

# array that stores the loss of each epoch
losses = []

# import pdb
# pdb.set_trace() # break point

# LEARNING ITERATIONS
for epoch in np.arange(self.n_epochs):
    #### ESTIMATION (FORWARD PASS)
    # X.shape == (n_samples, n_features)
    # w.shape == (n_features,)
    z = np.dot(X, w) + b # shape: (n_samples,)
    p_hat = self.__sigmoid(z)

    loss_epoch = self.__log_loss(y, p_hat)
    losses.append(loss_epoch)

    #### GRADIENT DESCENT UPDATES (BACKWARD PASS)
    # grad_w.shape: (n_features,)
    # grad_b: float
    grad_w, grad_b = self.__gradient(X, y, p_hat, w, self.alpha)
    w = w - self.learning_rate * grad_w # shape: (n_features)
    b = b - self.learning_rate * grad_b # float

    # pdb.set_trace()

```



```

if verbose and (epoch == 0 or (epoch + 1) % verbose == 0):
    print(f'[INFO] epoch={epoch + 1}/{self.n_epochs}, loss={loss_epoch:.7f}')

    ## code snippet to save the intermediate decision boundaries each `verbose` iterations
    w1, w2 = w
    x1_decision_line = np.array([X[:,0].min(), X[:,0].max()])
    x2_decision_line = -(b + (w1 * x1_decision_line)) / w2
    ax = sns.scatterplot(x=X[:, 0], y=X[:,1], hue=y)
    sns.lineplot(x=x1_decision_line, y=x2_decision_line, color='lightseagreen', ax=ax)
    ax.set_xlabel('x1')
    ax.set_ylabel('x2')
    ax.set_xlim(X[:,0].min(), X[:,0].max())
    ax.set_ylim(X[:,1].min(), X[:,1].max())
    ax.set_title('Decision Boundary on Training Samples')
    fig = ax.get_figure()
    fig.savefig(f'./plots/decision_boundary_epoch_{epoch+1:08d}.png')
    plt.close(fig) # to avoid showing the plot on jupyter

if verbose:
    losses = np.array(losses)
    print(f'\nFinal loss: {losses[-1]}')
    print(f'\nMean loss: {losses.mean()} +- {losses.std()}')

### ASSIGN THE TRAINED PARAMETERS TO THE PRIVATE ATTRIBUTES
self.__w = w
self.__b = b

def __predict_proba(self, X: ndarray, w, b) -> ndarray:
    '''Estimate the probability for the positive class of input samples.

    Parameters
    -----
    X: ndarray of shape (n_samples, n_features)
        Input samples.

    Returns
    -----
    ndarray of shape (n_samples,)
        The estimated probabilities for the positive class of input samples.
    ...

    assert self.is_fitted(), 'The instance is not fitted yet.'
    assert X.ndim == 2, f'X must b 2D. Passed: {X.ndim}'

```

```

z = np.dot(X, w) + b
p_hat = self.__sigmoid(z)

return p_hat

def predict_proba(self, X: ndarray) -> ndarray:
    '''Estimate the probability for the positive class of input samples.

    Parameters
    -----
    X: ndarray of shape (n_samples, n_features)
        Input samples.

    Returns
    -----
    ndarray of shape (n_samples,)
        The estimated probabilities for the positive class of input samples.
    ...

    assert self.is_fitted(), 'The instance is not fitted yet.'
    assert X.ndim == 2, f'X must b 2D. Passed: {X.ndim}'

    p_hat_dict = {}
    for i, class_name in self.__indexes_dict.items():
        p_hat = self.__predict_proba(X, self.__w_dict[i], self.__b_dict[i])
        p_hat_dict[i] = p_hat
    p_hat_df = pd.DataFrame(p_hat_dict)
    y_prob = np.array([max(p_hat_df.loc[idx, :]) for idx in p_hat_df.index], dtype=float)
    return y_prob

def predict(self, X: ndarray) -> ndarray:
    '''Predict the labels for input samples.

    Parameters
    -----
    X: ndarray of shape (n_samples, n_features)
        Input samples.

    Returns
    -----
    ndarray of shape (n_samples,)
        Predicted labels of input samples.
    ...

    assert self.is_fitted(), 'The instance is not fitted yet.'
    assert X.ndim == 2, f'X must b 2D. Passed: {X.ndim}'

```

```

p_hat_dict = {}
for i, class_name in self.__indexes_dict.items():
    p_hat = self.__predict_proba(X, self.__w_dict[i], self.__b_dict[i])
    p_hat_dict[i] = p_hat
p_hat_df = pd.DataFrame(p_hat_dict)
y_hat = np.array([self.classes_[np.argmax(p_hat_df.loc[idx, :])] for idx in p_hat_df.index], dtype=int)
return y_hat

```

Foi implementada a estratégia one vs rest. Nesta estratégia são treinados N classificadores binários, onde N é o número de classes presentes no dataset.

Essa abordagem requer que cada modelo preveja uma probabilidade de associação de classe ou uma pontuação semelhante à probabilidade. O argmax dessas pontuações (índice de classe com a maior pontuação) é então usado para prever uma classe.

```

In [8]: clf = LogisticRegression()
        clf

```

```

Out[8]: LogisticRegression()

```

```

In [9]: print(clf)

```

```

Learning rate: 0.001
Number of epochs: 1000
Regularization constant (alpha): 0.0001
Random state: 42

```

```

Trained?: False

```

Testing fit()

```

In [10]: # L2 regularization - alpha=100
         clf = LogisticRegression(alpha=100)
         clf.fit(X_train, y_train)
         clf.coef_

```

```

Out[10]: {0: array([0.00739051, 0.00881894]),
          1: array([ 0.02433   , -0.00676488]),

```

```
2: array([-0.00231138, -0.02670888]),  
3: array([-0.0065807 ,  0.00431483])}
```

Aqui temos coef_ como um dicionario que contém diversos vetores/arrays. Um para cada classe, onde esta classe é considerada a classe positiva (ou classe 1) e o resto a classe negativa (ou classe zero).

In [11]:

```
print(clf)
```

```
Learning rate: 0.001  
Number of epochs: 1000  
Regularization constant (alpha): 100  
Random state: 42
```

```
Trained?: True
```

In []:

Prediction

In [12]:

```
y_test_prob = clf.predict_proba(X_test)  
y_test_prob
```

```
Out[12]: array([0.47409219, 0.46242621, 0.50457997, 0.4737065 , 0.4678898 ,  
                0.46636472, 0.50846764, 0.46166893, 0.50403969, 0.46360868,  
                0.45820408, 0.47970154, 0.4668075 , 0.46513186, 0.46748508,  
                0.46045944, 0.4629722 , 0.47544673, 0.46542475, 0.46173071,  
                0.46221837, 0.45750718, 0.45772585, 0.49491591, 0.48858815,  
                0.48268906, 0.46694989, 0.46816814, 0.47194886, 0.45987544,  
                0.46074422, 0.476009 , 0.47500157, 0.46845686, 0.48492409,  
                0.49893926, 0.48387606, 0.4651031 , 0.45991636, 0.4927566 ,  
                0.45741569, 0.46247902, 0.4673067 , 0.49086304, 0.45963007,  
                0.46469914, 0.46708697, 0.45855165, 0.45651636, 0.46720025,  
                0.48034911, 0.45661335, 0.46161976, 0.4753206 , 0.46851203,  
                0.4861022 , 0.46948193, 0.48228624, 0.4672067 , 0.50608473,  
                0.49344277, 0.45927418, 0.46223231, 0.49537182, 0.50592074,  
                0.46117906, 0.46984897, 0.46204583, 0.46207431, 0.45760145,  
                0.45920577, 0.46410664, 0.46677068, 0.46597734, 0.49516557,  
                0.46751804, 0.46350326, 0.46690114, 0.46199413, 0.48142555,  
                0.46593696, 0.50355335, 0.4594675 , 0.48306316, 0.46525052,  
                0.46612099, 0.50080363, 0.45951899, 0.46421788, 0.49703229,  
                0.46338002, 0.46094773, 0.49014982, 0.46635592, 0.46457795,
```

```
0.46383817, 0.48344801, 0.45621677, 0.46298191, 0.50612341,
0.45731474, 0.4568881 , 0.461448 , 0.46133828, 0.50119464,
0.46548493, 0.47314353, 0.46040413, 0.46352773, 0.4882089 ,
0.47764357, 0.48115196, 0.45764654, 0.51078335, 0.51260291,
0.46201732, 0.46605829, 0.46539501, 0.48021379, 0.49125921,
0.49846957, 0.46289201, 0.4615041 , 0.47217623, 0.50762407,
0.45899456, 0.46047546, 0.46303314, 0.50710202, 0.46426484,
0.47154057, 0.49270745, 0.4635573 , 0.48451797, 0.46127366,
0.47458689, 0.45973765, 0.4647811 , 0.47590508, 0.47357102,
0.4921659 , 0.46323441, 0.5024931 , 0.48904736, 0.49823297,
0.48252643, 0.45908059, 0.46744035, 0.4655322 , 0.48708209,
0.50153918, 0.47836141, 0.46545001, 0.4778955 , 0.45625896,
0.46442082, 0.46403776, 0.46440845, 0.49572112, 0.47856832,
0.50196381, 0.46524966, 0.46526284, 0.50259494, 0.45766431,
0.46378045, 0.51123271, 0.4576839 , 0.50090846, 0.50349677,
0.48330628, 0.46157896, 0.46157931, 0.47979051, 0.48146498,
0.47801098, 0.48864623, 0.47692834, 0.46303559, 0.48964615,
0.45954665, 0.50929609, 0.49922404, 0.47459174, 0.50285373,
0.48772301, 0.46155808, 0.47043664, 0.51812476, 0.49430577,
0.45990061, 0.45761316, 0.48282736, 0.4654473 , 0.51031143,
0.46029477, 0.46325464, 0.48075829, 0.50930075, 0.47934231])
```

```
In [13]: y_test_pred = clf.predict(X_test)
y_test_pred
```

```
Out[13]: array([1, 3, 2, 1, 0, 3, 2, 0, 2, 0, 0, 1, 0, 0, 0, 3, 3, 2, 3, 3, 3, 0,
1, 2, 2, 2, 3, 0, 1, 0, 3, 3, 1, 1, 3, 1, 2, 1, 3, 0, 2, 0, 3, 3, 2,
0, 3, 1, 0, 0, 3, 1, 3, 1, 1, 3, 2, 1, 1, 3, 2, 2, 0, 3, 2, 2, 0,
0, 3, 3, 0, 0, 3, 3, 3, 2, 3, 3, 3, 3, 1, 0, 2, 3, 2, 3, 3, 2, 3,
3, 2, 3, 3, 1, 3, 3, 3, 1, 0, 3, 2, 0, 0, 3, 0, 2, 3, 1, 3, 3, 2,
1, 1, 0, 2, 2, 3, 0, 0, 1, 2, 2, 3, 0, 1, 2, 0, 0, 0, 2, 3, 1, 2,
3, 2, 0, 3, 0, 0, 1, 1, 1, 0, 2, 2, 2, 2, 0, 3, 3, 2, 2, 1, 3, 2,
3, 0, 3, 3, 2, 1, 2, 0, 3, 2, 0, 3, 2, 0, 2, 2, 2, 0, 3, 1, 1, 1,
1, 1, 3, 1, 0, 2, 2, 1, 2, 2, 3, 1, 2, 2, 0, 3, 1, 3, 2, 0, 3, 1,
2, 1])
```


```
In [14]: print(classification_report(y_test, y_test_pred))
```

	precision	recall	f1-score	support
0	0.98	0.90	0.94	49
1	1.00	0.98	0.99	41
2	1.00	1.00	1.00	53
3	0.92	1.00	0.96	57

accuracy			0.97	200
macro avg	0.97	0.97	0.97	200
weighted avg	0.97	0.97	0.97	200

```
In [15]: plot_confusion_matrix(clf, X_test, y_test, normalize='true')
```

```
Out[15]: <sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at 0x2d45bcfe4f0>
```



```
In [ ]:
```

Validating the implementation on Wine Dataset



```
In [16]: from sklearn.datasets import load_wine
         from sklearn.decomposition import PCA
         from sklearn.manifold import TSNE
```

```
In [17]: wine = load_wine()
         X = wine['data']
         y = wine['target']
```

```
In [18]: print(X.shape)
         print(y.shape)
         print(f'Labels: {np.unique(y)}')
```

```
(178, 13)
(178,)
Labels: [0 1 2]
```

```
In [19]: # splitting into train and test
         X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

         print(f'X_train.shape = {X_train.shape}')
```

```
print(f'y_train.shape = {y_train.shape}')
print(f'X_test.shape = {X_test.shape}')
print(f'y_test.shape = {y_test.shape}')
```

```
X_train.shape = (142, 13)
y_train.shape = (142,)
X_test.shape = (36, 13)
y_test.shape = (36,)
```

In []:

Data Exploration

In [20]:

```
df = pd.DataFrame(X_train, columns=wine['feature_names'])
df
```

Out[20]:

	alcohol	malic_acid	ash	alcalinity_of_ash	magnesium	total_phenols	flavanoids	nonflavanoid_phenols	proanthocyanins	color_intensity	hue	od280
0	14.34	1.68	2.70	25.0	98.0	2.80	1.31	0.53	2.70	13.00	0.57	
1	12.53	5.51	2.64	25.0	96.0	1.79	0.60	0.63	1.10	5.00	0.82	
2	12.37	1.07	2.10	18.5	88.0	3.52	3.75	0.24	1.95	4.50	1.04	
3	13.48	1.67	2.64	22.5	89.0	2.60	1.10	0.52	2.29	11.75	0.57	
4	13.07	1.50	2.10	15.5	98.0	2.40	2.64	0.28	1.37	3.70	1.18	
...	
137	13.86	1.51	2.67	25.0	86.0	2.95	2.86	0.21	1.87	3.38	1.36	
138	12.25	1.73	2.12	19.0	80.0	1.65	2.03	0.37	1.63	3.40	1.00	
139	14.38	1.87	2.38	12.0	102.0	3.30	3.64	0.29	2.96	7.50	1.20	
140	12.69	1.53	2.26	20.7	80.0	1.38	1.46	0.58	1.62	3.05	0.96	
141	12.34	2.45	2.46	21.0	98.0	2.56	2.11	0.34	1.31	2.80	0.80	

142 rows × 13 columns




```
In [21]: df.describe().transpose()
```

```
Out[21]:
```

	count	mean	std	min	25%	50%	75%	max
alcohol	142.0	12.979085	0.820116	11.03	12.3325	13.010	13.6775	14.83
malic_acid	142.0	2.373521	1.143934	0.89	1.6150	1.875	3.1350	5.80
ash	142.0	2.360845	0.279217	1.36	2.2100	2.360	2.5400	3.23
alcalinity_of_ash	142.0	19.473239	3.454792	10.60	17.2000	19.200	21.5000	30.00
magnesium	142.0	100.443662	14.650793	70.00	88.2500	98.000	107.0000	162.00
total_phenols	142.0	2.289085	0.637715	0.98	1.7250	2.310	2.8000	3.88
flavanoids	142.0	2.002113	1.004170	0.34	1.1250	2.075	2.8425	5.08
nonflavanoid_phenols	142.0	0.368028	0.128269	0.13	0.2700	0.340	0.4700	0.66
proanthocyanins	142.0	1.608028	0.583656	0.42	1.2500	1.555	1.9675	3.58
color_intensity	142.0	5.057606	2.330917	1.74	3.2200	4.600	6.1225	13.00
hue	142.0	0.956380	0.234101	0.48	0.7825	0.965	1.1200	1.71
od280/od315_of_diluted_wines	142.0	2.592817	0.722141	1.27	1.8375	2.775	3.1700	4.00
proline	142.0	734.894366	302.323595	278.00	502.5000	660.000	932.7500	1547.00

```
In [22]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 142 entries, 0 to 141
Data columns (total 13 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   alcohol                               142 non-null    float64
1   malic_acid                            142 non-null    float64
2   ash                                    142 non-null    float64
3   alcalinity_of_ash                     142 non-null    float64
4   magnesium                             142 non-null    float64
5   total_phenols                         142 non-null    float64
6   flavanoids                            142 non-null    float64
7   nonflavanoid_phenols                  142 non-null    float64
```

```
8  proanthocyanins          142 non-null    float64
9  color_intensity          142 non-null    float64
10 hue                      142 non-null    float64
11 od280/od315_of_diluted_wines 142 non-null    float64
12 proline                  142 non-null    float64
dtypes: float64(13)
memory usage: 14.5 KB
```

Preprocessing

```
In [23]: scaler = MinMaxScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)
```

```
In [ ]:
```


Components Analysis

```
In [24]: pca = PCA(n_components=3)
X_pca = pca.fit_transform(X_train)

tsne = TSNE(n_components=3)
X_tsne = tsne.fit_transform(X_train)
```

```
In [25]: plt.figure(figsize=(12, 6))
plt.subplot(1,2,1)
plt.title('PCA')
for y in np.unique(y_train):
    X_tmp = X_pca[y_train == y]
    plt.scatter(X_tmp[:, 0], X_tmp[:, 1], label=y, alpha=0.75)
plt.legend()
plt.subplot(1,2,2)
plt.title('TSNE')
for y in np.unique(y_train):
    X_tmp = X_tsne[y_train == y]
    plt.scatter(X_tmp[:, 0], X_tmp[:, 1], label=y, alpha=0.75)
plt.legend()
```

```
plt.show()
```



Pode ser observado que, ao aplicar o PCA, as três classes são distintas umas das outras, podendo ser separadas quase que perfeitamente. Já a técnica TSNE não apresenta a mesma capacidade quando considerando-se apenas as 2 componentes mais relevantes.

In []:

Baseline

In [26]:

```
clf = LogisticRegression()
clf.fit(X_train, y_train)
```

In [27]:

```
y_test_pred = clf.predict(X_test)
print(classification_report(y_test, y_test_pred))
```

	precision	recall	f1-score	support
0	1.00	0.93	0.96	14
1	0.61	1.00	0.76	14
2	0.00	0.00	0.00	8
accuracy			0.75	36
macro avg	0.54	0.64	0.57	36
weighted avg	0.63	0.75	0.67	36

C:\Users\weld1\Miniconda3\envs\work\lib\site-packages\sklearn\metrics_classification.py:1248: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.

```
_warn_prf(average, modifier, msg_start, len(result))
```

C:\Users\weld1\Miniconda3\envs\work\lib\site-packages\sklearn\metrics_classification.py:1248: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.


```
_warn_prf(average, modifier, msg_start, len(result))
```

C:\Users\weld1\Miniconda3\envs\work\lib\site-packages\sklearn\metrics_classification.py:1248: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.

```
_warn_prf(average, modifier, msg_start, len(result))
```

In [28]: `plot_confusion_matrix(clf, X_test, y_test, normalize='true')`

Out[28]: <sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at 0x2d48047aee0>



O modelo base apresentou uma performance bastante ruim. Este modelo foi treinado considerando os valores padrão da classe LogisticRegression. Na próxima seção será realizado o fine-tuning, visando obter melhores resultados no teste.

In []:

Fine-Tuning

In [29]:

```
n_jobs = 4
n_iter = 100
cv = 5
scoring = 'balanced_accuracy'
```

In [30]:

```
# LogisticRegression:
param_grid = {'learning_rate': [10, 1, 1e-1, 1e-2, 1e-3, 1e-4, 1e-5],
              'n_epochs': [1000, 5000, 1000, 2000],
              'alpha': [10, 1, 1e-1, 1e-2, 1e-3, 1e-4, 1e-5],
              }

logreg_clf = RandomizedSearchCV(LogisticRegression(), param_distributions=param_grid, n_iter=n_iter, cv=cv,
                               scoring=scoring, n_jobs=n_jobs, verbose=1, random_state=42)
logreg_clf.fit(X_train, y_train)
```

Fitting 5 folds for each of 100 candidates, totalling 500 fits

Out[30]:

```
RandomizedSearchCV(cv=5, estimator=LogisticRegression(), n_iter=100, n_jobs=4,
                  param_distributions={'alpha': [10, 1, 0.1, 0.01, 0.001,
                                                0.0001, 1e-05],
                  'learning_rate': [10, 1, 0.1, 0.01,
                                    0.001, 0.0001,
                                    1e-05],
                  'n_epochs': [1000, 5000, 1000, 2000]},
                  random_state=42, scoring='balanced_accuracy', verbose=1)
```

In [31]:

```
model = logreg_clf.best_estimator_
model.fit(X_train, y_train)
```


In [32]:

```
y_test_pred = model.predict(X_test)
print(classification_report(y_test, y_test_pred))
```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	14
1	1.00	1.00	1.00	14
2	1.00	1.00	1.00	8
accuracy			1.00	36
macro avg	1.00	1.00	1.00	36
weighted avg	1.00	1.00	1.00	36

In [33]: `plot_confusion_matrix(model, X_test, y_test, normalize='true')`

Out[33]: `<sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at 0x2d4571adca0>`



Após o fine-tuning pudemos notar que o modelo obteve uma melhoria significativa em relação ao baseline. Este é um dataset relativamente simples e pequeno. Não era esperado um resultado perfeito na base de teste.

In []:

Conclusion

Pudemos observar que a classe foi estendida para que pudesse lidar com problemas multiclasse, para isso foi escolhida a estratégia one vs all, devido

a sua simplicidade e se adequar melhor para problemas de muitas classes, se comparado à estratégia one vs one.

Considerando o dataset de validação, wine dataset, tivemos um baseline bastante ruim, porém ao fazer o fine-tuning obtivemos um excelente resultado.

Não foram implementadas diferentes versões do gradiente descendente.

In []: